

06 - Memory Management

January 23, 2023

COMP2404

Darryl Hill

Contents

1. Stack and Heap
2. Pointers
3. Memory allocation
4. Dynamic arrays

OS allocates 4 areas of memory on program start-up

- ▶ ***code segment*** or ***text segment***
 - ▶ Stores the program instructions
- ▶ ***data segment***
 - ▶ Contains global memory
 - ▶ global data members, string literals, etc
- ▶ ***function call stack***
 - ▶ Variables and data associated with current function
 - ▶ ***Stack frames***
- ▶ ***heap*** - part of the data segment
 - ▶ Dynamically allocated memory

The ***Code Segment*** contains

- ▶ program instructions
- ▶ function addresses

The ***Data Segment*** contains

- ▶ global variables
- ▶ static variables
- ▶ literals

The *Function Call Stack*:

- ▶ keeps track of functions called and the order in which they were called
- ▶ stores local variables
- ▶ stores function parameters
- ▶ function meta-data


The *Heap*:

- ▶ is part of the data segment
- ▶ stores all dynamically allocated memory
 - ▶ memory the programmer manages
 - ▶ more on this soon

Function Call Stack

A **Stack** is a type data structure

- ▶ collection of related data
 - ▶ last in, first out (LIFO)
 - ▶ similar to a pile of dishes
-
- ▶ adding an item is referred to as **pushing** the item
 - ▶ removing an item is referred to as **popping** the item



The **Function Call Stack** is a data structure to keep track of function calls.

- ▶ When a function is called, a **stack frame** is pushed onto the call stack.
- ▶ When a function returns, the **stack frame** is popped off.

A **stack frame** contains:

- ▶ Local variables contained in the function parameters.
- ▶ The return address of the next instruction in the calling function.

The **Heap** is an area of memory used for **dynamic allocation** (more on this next slide).

- ▶ Allocation that happens at runtime, using the **new** keyword.
- ▶ This is memory that is managed by the programmer.
 - ▶ We must **delete** it at some point.



Figure: "The heaping pile o' mail again" by Charles Williams is licensed under CC BY 2.0

Statically allocated memory:

- ▶ Memory is reserved **where the variable was declared**.
 - ▶ Either on the function call stack or in a class.
 - ▶ Variables (of any type) are simply declared
 - ▶ `int x; Student stu;`
- ▶ This memory is deleted for us when the containing object is *deleted* or stack frame is popped off.

Dynamically allocated memory:

- ▶ Memory is reserved **on the heap**.
 - ▶ Uses the `new` keyword which returns a pointer (memory address from the heap).
 - ▶ `int* x = new int; Student* stu = new Student;`
- ▶ We must `delete` this memory explicitly.

Pointers

A ***Pointer*** is a variable that stores a ***memory address***.

- ▶ The address of a variable, or
- ▶ the start of a block of dynamically allocated memory (an array).

Why use pointers?

- ▶ They have a small, fixed size – 8 bytes.
 - ▶ Often smaller than the data it points to.
 - ▶ Easier to pass (the pointer) by value.
- ▶ They are an easy way to avoid copying data.
 - ▶ C++ can make many needless copies.
- ▶ Pointers are ***necessary*** for dynamically allocated memory.
 - ▶ `new` returns a pointer.

Like all variables, pointers have a:

- ▶ Name
- ▶ Data type – the type of data being pointed to.
 - ▶ Tells the compiler how to interpret the data found at that memory location.
- ▶ Value
 - ▶ This value is an address in memory.
 - ▶ This address must contain a variable of the type associated with the pointer.
- ▶ Location in memory
 - ▶ Pointers have their own memory addresses.

Memory Map

Name	Type	Value	Location



A8: char c



A9: int i



AA: char* cptr

- ▶ In 64-bit architecture, pointers are 64-bits (8 bytes) of memory.
- ▶ Pointers may point to a variable of
 - ▶ any data type (but each pointer is restricted to one type), and
 - ▶ any size.
- ▶ Thus we can “pass” a variable into a function by copying 8 bytes of data, irrespective of the size of the variable.

Coding example <p1>

Variable declaration is a statement specifying variable's name and data type.

A pointer variable declaration has 3 components:

- ▶ The type of data being pointed to,
- ▶ the * symbol to indicate that it is a pointer, and
- ▶ the variable name.

Pointers can receive a value of

- ▶ the memory address of a variable, or
- ▶ a system call for a new block of memory.

The & symbol is the “address-of” operator

- ▶ It returns the memory location of the operand variable.

```
int b;  
int* a = &b;  
int* c = new int;
```

The * symbol:

- ▶ is the ***dereferencing*** operator
- ▶ returns the value stored at the memory address pointed to by the operand

NULL pointer exception:

- ▶ this happens when dereferencing a pointer that is set to **NULL**.
 - ▶ the pointer value is 0, so we try and access memory address 0

Good programming practice to check for **NULL** pointers

- ▶ same as in Java
- ▶ also good practice: set garbage pointers to **NULL** so that you may **NULL**-check them

Symbols * and & have multiple roles

In variable declarations:

- ▶ * indicates a pointer
- ▶ & indicates a reference

Anywhere else

- ▶ * is dereferencing
- ▶ & is address-of

Both have another meaning as a math and bit operator

Both are overloaded symbols, so pay attention to context!

Parameter Passing

Pass-by-value:

- ▶ this makes copies of the values of the variables
- ▶ can be very inefficient (though it is ok for primitives)
- ▶ does not use the variable address

Pass-by-reference has 2 forms:

- ▶ using references
- ▶ using pointers

Java does pass-by-reference, except for primitives which are pass by value

- ▶ Programmer has no choice in this

Coding example <p2>

Memory allocation:

- ▶ reserving a specific number of bytes in memory based on data type

Static Memory Allocation

- ▶ done by declaring a variable
- ▶ memory reserved wherever it is declared
- ▶ memory is deallocated when the stack frame that contains it is popped off or containing object is deallocated

Dynamic Memory Allocation

- ▶ memory is reserved on the heap using the `new` keyword
- ▶ memory is deallocated by using `delete` keyword

coding example <p3>

Memory Leaks

C/C++ does not provide garbage collection

- ▶ we can have memory leaks

What is a memory leak?

- ▶ a block of dynamically allocated memory with no pointers to it
- ▶ can never again be used by the program (and while the program is running, no other programs can use it either)
- ▶ cannot call `delete` on it

What causes memory leaks?

- ▶ a pointer gets **clobbered** – overwritten
- ▶ a pointer moves out of scope

Memory Leaks

Why are memory leaks a problem?

- ▶ access to data permanently lost
 - ▶ finite amount of heap space given to the program

Dynamically allocated memory is released

- ▶ once the program terminates
- ▶ explicitly released with the `delete` keyword

Industrial grade software should run for months or years

- ▶ if code repeats (and most code does), the memory leak will repeat
- ▶ could run out of heap space

coding example `<p3>` **redux**

Preventing Memory Leaks

Always explicitly deallocate memory

- ▶ for each **new** there should be a matching **delete**
- ▶ use valgrind utility in Linux to check

“Where is this object deleted?” is something you should ask yourself again and again.

Good rule of thumb: objects are deleted where they are created

HOWEVER

- ▶ occasionally created in one place and deleted in another
- ▶ be very careful doing this
- ▶ **Factory** design pattern is one example
- ▶ Often requires tight coupling between classes or functions
- ▶ Advanced design



Every dynamically allocated object must be **deleted** in **exactly one place**

- ▶ Often deleted by the containing object (not a data structure) or the declaring function.
- ▶ Containers often form a “hierarchy of deletion responsibility”



Also, make sure you don't clobber a pointer into the heap

- ▶ Each piece of data should have a primary reference.
 - ▶ This reference should exist for the lifetime of the object
- ▶ No garbage collection or reference counting in this class
- ▶ Even with these “memory helpers”, if you do not understand memory fundamentals, garbage collection or deallocation can occur at the wrong time
 - ▶ This will cause a slow-down

Dynamic Allocation of Objects

Using `new` calls the object **constructor**

- ▶ `malloc` does not call the constructor- DO NOT USE
- ▶ parameters to the constructor can be specified
- ▶ memory is allocated on the heap
- ▶ member variables initialized

Using `delete` calls the object **destructor**

- ▶ `free` does not call the destructor - DO NOT USE
- ▶ resources are cleaned up
- ▶ memory is released

coding example <p4>

Returning Objects by Reference - Double Pointers

Sometimes we want to return an object by reference

- ▶ Using return parameters
- ▶ Often from the heap, but can be from a (previous) stack frame
- ▶ This parameter should return the *address* of this object
- ▶ Usually in the form of a *pointer*

To do this we pass *in* a *pointer* by *reference*.

Double Pointers

Memory Map

Name	Type	Value	Location



A8: int i



AC: int* iptr



B3: int** dptr

Retrieving Dynamic Objects - Double Pointers

Double pointer syntax: (coding example <p5>)

```
int main(){
    Date* date1;
    getDate(&date1);
    date1->print();    \\ happy new year!
}

\\ this will return a pointer
\\ to a dynamically allocated object
void getDate(Date** d){
    Date* date2 = new Date(2022,1,1);
    *d = date2;
}
```

Do **NOT** declare a double pointer.
Declare a pointer and pass in the address.

Memory Map


Name	Type	Value	Location



A4: Date* date1



AC: Date* date2



DC: Date object



B3: Date** d

Dynamically Allocated Arrays

Arrays are dynamically allocated with the `new[]` operator

- ▶ if you make an array of objects, the default constructor is called for each of these objects
 - ▶ if it exists
 - ▶ if not, you will get an error

Dynamic arrays can be deallocated with the `delete[]` operator

- ▶ if it is an array of objects, the destructor is called for every element of the array

Using `delete` without brackets:

- ▶ may not compile
- ▶ might result in undefined behaviour
- ▶ might work `~_(ツ)_/~`

Arrays also have two types of memory allocation:

- ▶ static
- ▶ dynamic

In addition, arrays can store

- ▶ objects or
- ▶ object pointers

That means there are 4 ways of storing data in arrays

coding example <p6>

Statically allocated array of objects:

- ▶ `Date dates[size];`
- ▶ `dates[0].print();`

Dynamically allocated array of objects:

- ▶ `Date* dates = new Date[size];`
- ▶ `dates[0].print();`

Statically allocated array of pointers:

- ▶ `Date* dates[size];`
- ▶ `dates[0] = new Date;`
- ▶ `dates[0]->print();`

Dynamically allocated array of pointers:

- ▶ `Date** dates = new Date*[size];`
- ▶ `dates[0] = new Date;`
- ▶ `dates[0]->print();`